

Report on CAN-PR2

Submitted by Cyril WU, 5132609029, Hugo BI, 5132609036, Nicolas YING, 5030369019.

3 Pipelining

The CPU supports forwarding.

e.g. Cycle 3, instruction 0x11d8: sw 28(r29), r31 is in EX stage, where r29 isn't yet written back (in MEM stage).

| Instructions/Cycles | 0 | 1 | 2 | 3 | 4 | 5 |
|---------------------|----|----|----|-----|-----|-----|
| addui r29,r29,-32 | IF | ID | EX | MEM | WB | |
| sw 28(r29),r31 | | IF | ID | EX | MEM | WB |
| sw 24(r29),r30 | | | IF | ID | EX | MEM |

```
DEBUG [openDLX]: -----
DEBUG [openDLX]: Cycle 3 start
DEBUG [openDLX]: -----
DEBUG [ICache]: idx: 1f index mask: 0x1f value: 0x000011e0 result: 0x1c
DEBUG [ICache]: Accessing index: 28
DEBUG [ICache]: idx: 1f index mask: 0x1f value: 0x000011e0 result: 0x1c
DEBUG [ICache]: Accessing index: 28
DEBUG [ICache]: idx: 1f index mask: 0x1f value: 0x000011e0 result: 0x1c
DEBUG [ICache]: Accessing way: 0 index: 28
DEBUG [FETCH]: PC: 0x000011e0 fetched instruction 0x03a0f025
DEBUG [DECODE]: PC: 0x000011dc instruction decoded as SW ALU:ADD rRs:true (29/sp) rRt:true (30/s8) wRt:fa
DEBUG [EXECUTE]: {FW} using ALU result 0xffffffe0 for register 29/sp instead of value: 0x00000000
DEBUG [EXECUTE]: {FW} PC: 0x000011d8 forwarding changed value for ALU port A RS from: 0x0 to: 0xffffffe0
DEBUG [EXECUTE/ALU]: -32(0xffffffe0) + 28(0x0000001c) = -4(0xfffffff4)
DEBUG [EXECUTE]: PC: 0x000011d8 ALU calculated: -4(0xfffffff4) by: -32(0xffffffe0) ADD 28(0x0000001c)
DEBUG [BP_MODULE]: instruction at: 0x000011dc was not found in BTB
DEBUG [MEMORY]: PC: 0x000011d4 nothing to do
INFO [WRITEBACK]: PC: 0x00000000
DEBUG [FETCH]: PC: 0x000011e0 fetched instruction 0x03a0f025
```

In addition, it is explicitly shown from the log that a forwarding is done to fetch the computed r29 value for sw instruction.

Special case for lw

When 0x11ec lw is in EX and MEM stages, slti needs r2 value to enter EX stage, thus this dependency caused a data hazard, solved by a stalling of penalty 1.

| | | | | | | |
|---------------|----|----|----|-----|-----|----|
| lhi r2,1 | IF | ID | EX | MEM | WB | |
| lw r2,24(r30) | | IF | ID | EX | MEM | WB |
| slti r2,r2,8 | | | IF | ID | ID | EX |


Stages striked out

| | | | | | | | |
|------------|----|----|----|-----|-----|-----|------------|
| jal 0x1098 | IF | ID | EX | MEM | WB | | 0x000011e8 |
| nop | | IF | ID | EX | MEM | WB | 0x000011ec |
| nop | | | IF | ID | EX | MEM | WB |

0x11f0 is actually a misprediction produced at 0x11e8, (it is not actually a prediction since the instuction

0x11e8 is an unconditional jump). Thus when 0x11e8 enters EX stage and the CPU spots out that 0x11f0 should not be executed at all, thus it is flushed and the correct instruction is fetched from 0x1098.

Another example could be find here:



and here:



The latter two was produced by misprediction of conditional branches.

Difference of the methodology of handling jumps and branches

1. The lecture introduced a processor with dedicated branch prediction, selection and address calculation circuit, while the mocked processor integrates these tasks into ALU.
2. Our binary file have always at least one nop following a (conditional or unconditional) jump, thus we only have to flush one instruction in the case of one misprediction rather than two instructions. This feature could be implemented by the compiler.

4. Branch Prediction

Comparison of two predictors

We compare the performance of 1 bit predictor and 2 bit predictor:

log file of 1 bit predictor

INFO [openDLX]: Jumps correctly predicted: 88 mispredicted: 37 misprediction rate: 29,6%

log file of 2 bit predictor

INFO [openDLX]: Jumps correctly predicted: 99 mispredicted: 26 misprediction rate: 20,8%

We can see that the number of predictions are both 125 and the 1 bit prediction approach has a higher misprediction rate.

Detail

1 bit predictor

INFO [openDLX]: bpc: 0x00001078 [120] tgts: [0x00001024] a:44 t/nt: 36/8 mp/cp: 16/28 mp-ratio: 0,36

```

INFO  [openDLX]: bpc: 0x00001050 [80] tgts: [0x00001060] a:36 t/nt: 31/5 mp/cp: 11/25 mp-
ratio: 0,31
INFO  [openDLX]: bpc: 0x000010f4 [116] tgts: [0x000010b4] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio:
0,22
INFO  [openDLX]: bpc: 0x000011ac [44] tgts: [0x00001108] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio:
0,22
INFO  [openDLX]: bpc: 0x0000101c [28] tgts: [0x0000106c] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x00001090 [16] tgts: [0x00001130] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x00001128 [40] tgts: [0x00001000] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x000010ac [44] tgts: [0x000010ec] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio:
1
INFO  [openDLX]: bpc: 0x00001100 [0] tgts: [0x000011a4] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1
INFO  [openDLX]: bpc: 0x000011e8 [104] tgts: [0x00001098] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio:
1

```

2 bit predictor

```

INFO  [openDLX]: bpc: 0x00001078 [120] tgts: [0x00001024] a:44 t/nt: 36/8 mp/cp: 9/35 mp-
ratio: 0,2
INFO  [openDLX]: bpc: 0x00001050 [80] tgts: [0x00001060] a:36 t/nt: 31/5 mp/cp: 7/29 mp-
ratio: 0,19
INFO  [openDLX]: bpc: 0x000010f4 [116] tgts: [0x000010b4] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio:
0,22
INFO  [openDLX]: bpc: 0x000011ac [44] tgts: [0x00001108] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio:
0,22
INFO  [openDLX]: bpc: 0x0000101c [28] tgts: [0x0000106c] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x00001090 [16] tgts: [0x00001130] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x00001128 [40] tgts: [0x00001000] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio:
0,12
INFO  [openDLX]: bpc: 0x000010ac [44] tgts: [0x000010ec] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio:
1
INFO  [openDLX]: bpc: 0x00001100 [0] tgts: [0x000011a4] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1
INFO  [openDLX]: bpc: 0x000011e8 [104] tgts: [0x00001098] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio:
1

```

When we look at the situation of each branch, we find that the difference between two predictors is more significant when the branch is executed lots of times. For example, the branch **0x00001024**. However, 2 bit predictor is not always better than 1 bit predictor. On branches **0x000010b4** to **0x00001098**, performances of two predictors are the same.

Explanation of the difference

1078: 1440ffea bnez v0,1024 <minIndex+0x24>

C Code **<minIndex>**: `for(i = 0; i < n; i++)`, the condition after for-loop. It is usually taken.

1050: 10400003 beqz v0,1060 <minIndex+0x60>

C Code <minIndex>: if (array[i] < array[minIdx]), usually taken.

10f4: 1440ffef bnez v0,10b4 <main+0x1c>

C Code <main>: for(i = 0; i < SIZE; i++), the condition after for-loop. It is usually taken.

11ac: 1440ffd6 bnez v0,1108 <main+0x70>

C Code <main>: for(i = 0; i < SIZE; i++), the condition after for-loop. It is usually taken.

101c: 10000013 b 106c <minIndex+0x6c>

C code <minIndex>: for(i = 0; i < n; i++), the condition before for-loop. It is always taken and PC jumps to the condition after for-loop.

1090: 03e00008 jr ra

C Code <minIndex>: return minIdx, Always taken.

1128: 0c000400 jal 1000 <minIndex>

C Code <main>: minIndex(&buf[i], SIZE - i), Always taken.

10ac: 1000000f b 10ec <main+0x54>

1100: 10000028 b 11a4 <main+0x10c>

C Code <main>: the condition before for-loop. It is usually taken.

11e8: 0c000426 jal 1098 <main>

C Code <__start>: main(), Always taken.

The 1 bit predictor do its prediction based on the recent result of the branch. We firstly assume that each branch has its own bit, i.e for every branch instruction address, a mod k are not equal to other (no collision in hash table).

The instruction 1078 correspond to the condition of exiting the for-loop in minIndex function. Function minIndex is called 8 times in the programme. The 1 bit predictor will have a misprediction in the first loop and the last loop, thus we have 16 mispresictions. Meanwhile, the 2 bit predictor, at the beginning, has the 2-bit 00. In the first two loops, the predictor will not take the branch, because the bit patterns are 00 and 01. After that, the predictor will take the branch because of its bit partterns: 11 or 10 and we can see that the bit pattern will be always 11 and 10. So 2 bit predictor will only has one misprediction when exiting the loop, added with the first two mispredictions, it has in total 9 mispredictions.

Penalties

We find that there is always a nop instruction after a conditional jump instruction, which is, in our opinion, systematically added by the compiler, and the penalties (including the nop added) are:

PC/Prediction taken/untaken: 3

PC/Prediction taken/taken: 2

PC/Prediction untaken/taken: 3


PC/Prediction untaken/untaken: 2

Because there is a nop after condition, CPU only need to flush when a misprediction occur. In addition, CPU has the same behavior after a taken/untaken or a untaken/taken, which could simplify the design.

5 Data Caches

memory_latency

We give a delay time between the memory's receipt of a read request and its release of data corresponding with the request. So the IF, MEM and WB will have delay in the pipelining. Other stages will also have delays for the data dependency.



Dcache:

Accesses: 500 hits: 220 misses: 280 hit rate: 44% total size: 16

dcache_total_block_number

dcache_total_block_number = 4

Dcache:

total size: 16

Accesses: 500 hits: 220 misses: 280 hit rate: 44%

dcache_total_block_number = 8

Dcache:

total size: 32

Accesses: 499 hits: 381 misses: 118 hit rate: 76.35%

dcache_total_block_number = 16

Dcache:

total size: 64

Accesses: 500 hits: 411 misses: 89 hit rate: 82.2%

dcache_total_block_number = 32

Dcache:

total size: 128

Accesses: 500 hits: 492 misses: 8 hit rate: 98.4%

dcache_total_block_number = 64

Dcache:

total size: 256

Accesses: 500 hits: 492 misses: 8 hit rate: 98.4%

Conclusion:

The cache size is changed from 16 to 256, and the hit rate has a continuous growth. With regard to the cache

size, the case when `dcache_total_block_number = 32` performs best.

`dcache_associativity` with `dcache_total_block_number = 8`

`dcache_associativity = 2`, `dacache_replacement_policy = FIFO`

```
Dcache:  
hit rate: 74.55%
```

`dcache_associativity = 2`, `dacache_replacement_policy = LRU`

```
Dcache:  
hit rate: 77.35%
```

`dcache_associativity = 4`, `dacache_replacement_policy = FIFO`

```
Dcache:  
total size: 32  
rate: 83.97%
```

`dcache_associativity = 4`, `dacache_replacement_policy = LRU`

```
Dcache:  
total size: 32  
hit rate: 85.97%
```

`dcache_associativity = 8`, `dacache_replacement_policy = FIFO`

```
Dcache:  
total size: 32  
hit rate: 83.37%
```

`dcache_associativity = 8`, `dacache_replacement_policy = LRU`

```
Dcache:  
total size: 32  
hit rate: 87.58%
```

`dcache_associativity = 1`, `dacache_replacement_policy = DIRECT_MAPPED`

```
Dcache:  
total size: 32  
hit rate: 76.35%
```

Conclusion:

The two replacement policies with a higher data cache associativity have better performances than the original `DIRECT_MAPPED` with the same total size in general.

dcache_block_size with LRU 4 way associativity and dcache_total_block_number = 8

dcache_block_size = 4

```
total size: 32  
hit rate: 85.97%
```

dcache_block_size = 8

```
total size: 64  
hit rate: 96.8%
```

dcache_block_size = 16

```
total size: 128  
hit rate: 99.6%
```

dcache_block_size = 32

```
total size: 256  
hit rate: 99.6%
```

Conclusion:

Compared with block number section, the scheme to increase the block size has better performance than the scheme to increase the block number, with the same total cache size.

C code examination

The data are accessed by the program in array buf and in counters (spatial locality and temporal locality), so the data is reused frequently.

The capacity miss dominates the cache's performance in this program because the data used is very fixed.

The program actually access the data:

1. in main, SIZE numbers to sort: buf[SIZE]
2. in main, counter i
3. in main, the temporary variable: tmp and minIdx
4. in minIndex, the local variable: i and minIdx

in total $8+1+2+2 = 13$.

Practically, a cache that is large enough to contain 8 integers (32 cache size units) would provide a promising hit rate as they are the most frequently used data. In addition, if a data cache is large enough to cache all the data accessed, data miss would only be generated the first time it is accessed.

Final question

The configuration `dcache_block_size = 16/32`, 4-way set-associative, total number of cache blocks = 8 performs the best (the hit rate 99.6% exceeds all the others cases).

The configuration: `dcache_block_size = 8`, 4-way set-associative, total number of cache blocks = 8 made the best trade-off between the total cache size and the hit rate. Because the hit rate is rather good and close to 100% (96.8%), but the total cache size is merely half of the best performance case (the case total size = 128, hit rate = 99.6%) in block size section.

From the result of changing data cache number and size, we can observe that as the total size increase, the hit rate is significantly enhanced, and from question 4 that as the total cache size is fixed, the enhance of hit rate is limited. So we can conclude that the total cache size is the most relevant parameter to the hit rate. And this conclusion is coherent with the analysis of data access that the data is reused very frequently and the most dominant factor of the hit rate is the cache size.

2 Warm Up

2.1 MIPS Tools

Compile insertion-sort.c

```
bin/insertion-sort.elf: src/insertion-sort.c
    $(MIPS_CC) $(MIPS_CFLAGS) -o $@ $<

bin/insertion-sort.bin: bin/insertion-sort.elf
    $(MIPS_OBJCOPY) -O binary -S $< $@

cfg/*.cfg: bin/insertion-sort.elf
    bin/updatecfgFiles.sh $< ./cfg > /dev/null 2>&1
```

-nostdlib and -nostartfiles

```
-nostdlib
    Do not use the standard system startup files or libraries when linking. No
    startup files and only the
        libraries you specify will be passed to the linker, options specifying linkage of
    the system libraries,
        such as "--static-libgcc" or "--shared-libgcc", will be ignored. The compiler may
    generate calls to
        "memcmp", "memset", "memcpy" and "memmove". These entries are usually resolved by
    entries in libc.
        These entry points should be supplied through some other mechanism when this
    option is specified.

    One of the standard libraries bypassed by -nostdlib and -nodefaultlibs is
    libgcc.a, a library of
        internal subroutines which GCC uses to overcome shortcomings of particular
    machines, or special needs
        for some languages.
```


In most cases, you need `libgcc.a` even when you want to avoid other standard libraries. In other words, when you specify `-nostdlib` or `-nodefaultlibs` you should usually specify `-lgcc` as well. This ensures that you have no unresolved references to internal GCC library subroutines. (For example, `__main`, used to ensure C++ constructors will be called.)

`-nostdlib`

Do not use the standard system startup files or libraries when linking. No startup files and only the libraries you specify will be passed to the linker, options specifying linkage of the system libraries, such as `"-static-libgcc"` or `"-shared-libgcc"`, will be ignored. The compiler may generate calls to `"memcmp"`, `"memset"`, `"memcpy"` and `"memmove"`. These entries are usually resolved by entries in `libc`. These entry points should be supplied through some other mechanism when this option is specified.

One of the standard libraries bypassed by `-nostdlib` and `-nodefaultlibs` is `libgcc.a`, a library of internal subroutines which GCC uses to overcome shortcomings of particular machines, or special needs for some languages.

In most cases, you need `libgcc.a` even when you want to avoid other standard libraries. In other words, when you specify `-nostdlib` or `-nodefaultlibs` you should usually specify `-lgcc` as well. This ensures that you have no unresolved references to internal GCC library subroutines. (For example, `__main`, used to ensure C++ constructors will be called.)

look at the assembly code

```
sjtu@sjtu-OptiPlex-3010:~/Documents/DossierNicolasien/PR2$ mips-linux-gnu-objdump -d bin/insertion-sort.elf
```

```
bin/insertion-sort.elf:      file format elf32-tradlittlemips
```

Disassembly of section `.text`:

```
00001000 <minIndex>:
 1000:      27bdffe8      addiu   sp,sp,-24
 1004:      afbe0014      sw      s8,20(sp)
 1008:      03a0f025      move    s8,sp
 100c:      afc40018      sw      a0,24(s8)
```

| | | | |
|-------|----------|-------|-------------------------|
| 1010: | afc5001c | sw | a1,28(s8) |
| 1014: | afc0000c | sw | zero,12(s8) |
| 1018: | afc00008 | sw | zero,8(s8) |
| 101c: | 10000013 | b | 106c <minIndex+0x6c> |
| 1020: | 00000000 | nop | |
| 1024: | 8fc20008 | lw | v0,8(s8) |
| 1028: | 00021080 | sll | v0,v0,0x2 |
| 102c: | 8fc30018 | lw | v1,24(s8) |
| 1030: | 00621021 | addu | v0,v1,v0 |
| 1034: | 8c430000 | lw | v1,0(v0) |
| 1038: | 8fc2000c | lw | v0,12(s8) |
| 103c: | 00021080 | sll | v0,v0,0x2 |
| 1040: | 8fc40018 | lw | a0,24(s8) |
| 1044: | 00821021 | addu | v0,a0,v0 |
| 1048: | 8c420000 | lw | v0,0(v0) |
| 104c: | 0062102a | slt | v0,v1,v0 |
| 1050: | 10400003 | beqz | v0,1060 <minIndex+0x60> |
| 1054: | 00000000 | nop | |
| 1058: | 8fc20008 | lw | v0,8(s8) |
| 105c: | afc2000c | sw | v0,12(s8) |
| 1060: | 8fc20008 | lw | v0,8(s8) |
| 1064: | 24420001 | addiu | v0,v0,1 |
| 1068: | afc20008 | sw | v0,8(s8) |
| 106c: | 8fc30008 | lw | v1,8(s8) |
| 1070: | 8fc2001c | lw | v0,28(s8) |
| 1074: | 0062102a | slt | v0,v1,v0 |
| 1078: | 1440ffea | bnez | v0,1024 <minIndex+0x24> |
| 107c: | 00000000 | nop | |
| 1080: | 8fc2000c | lw | v0,12(s8) |
| 1084: | 03c0e825 | move | sp,s8 |
| 1088: | 8fbe0014 | lw | s8,20(sp) |
| 108c: | 27bd0018 | addiu | sp,sp,24 |
| 1090: | 03e00008 | jr | ra |
| 1094: | 00000000 | nop | |

00001098 <main>:

| | | | |
|-------|----------|-------|------------------|
| 1098: | 27bdffb0 | addiu | sp,sp,-80 |
| 109c: | afbf004c | sw | ra,76(sp) |
| 10a0: | afbe0048 | sw | s8,72(sp) |
| 10a4: | 03a0f025 | move | s8,sp |
| 10a8: | afc00018 | sw | zero,24(s8) |
| 10ac: | 1000000f | b | 10ec <main+0x54> |
| 10b0: | 00000000 | nop | |
| 10b4: | 3c020001 | lui | v0,0x1 |
| 10b8: | 8fc30018 | lw | v1,24(s8) |
| 10bc: | 00031880 | sll | v1,v1,0x2 |
| 10c0: | 24421210 | addiu | v0,v0,4624 |
| 10c4: | 00621021 | addu | v0,v1,v0 |
| 10c8: | 8c430000 | lw | v1,0(v0) |

| | | | |
|-------|----------|-------|---------------------|
| 10cc: | 8fc20018 | lw | v0,24(s8) |
| 10d0: | 00021080 | sll | v0,v0,0x2 |
| 10d4: | 27c40018 | addiu | a0,s8,24 |
| 10d8: | 00821021 | addu | v0,a0,v0 |
| 10dc: | ac43000c | sw | v1,12(v0) |
| 10e0: | 8fc20018 | lw | v0,24(s8) |
| 10e4: | 24420001 | addiu | v0,v0,1 |
| 10e8: | afc20018 | sw | v0,24(s8) |
| 10ec: | 8fc20018 | lw | v0,24(s8) |
| 10f0: | 28420008 | slti | v0,v0,8 |
| 10f4: | 1440ffef | bnez | v0,10b4 <main+0x1c> |
| 10f8: | 00000000 | nop | |
| 10fc: | afc00018 | sw | zero,24(s8) |
| 1100: | 10000028 | b | 11a4 <main+0x10c> |
| 1104: | 00000000 | nop | |
| 1108: | 27c30024 | addiu | v1,s8,36 |
| 110c: | 8fc20018 | lw | v0,24(s8) |
| 1110: | 00021080 | sll | v0,v0,0x2 |
| 1114: | 00622021 | addu | a0,v1,v0 |
| 1118: | 24030008 | li | v1,8 |
| 111c: | 8fc20018 | lw | v0,24(s8) |
| 1120: | 00621023 | subu | v0,v1,v0 |
| 1124: | 00402825 | move | a1,v0 |
| 1128: | 0c000400 | jal | 1000 <minIndex> |
| 112c: | 00000000 | nop | |
| 1130: | 00401825 | move | v1,v0 |
| 1134: | 8fc20018 | lw | v0,24(s8) |
| 1138: | 00621021 | addu | v0,v1,v0 |
| 113c: | afc2001c | sw | v0,28(s8) |
| 1140: | 8fc2001c | lw | v0,28(s8) |
| 1144: | 00021080 | sll | v0,v0,0x2 |
| 1148: | 27c30018 | addiu | v1,s8,24 |
| 114c: | 00621021 | addu | v0,v1,v0 |
| 1150: | 8c42000c | lw | v0,12(v0) |
| 1154: | afc20020 | sw | v0,32(s8) |
| 1158: | 8fc20018 | lw | v0,24(s8) |
| 115c: | 00021080 | sll | v0,v0,0x2 |
| 1160: | 27c30018 | addiu | v1,s8,24 |
| 1164: | 00621021 | addu | v0,v1,v0 |
| 1168: | 8c43000c | lw | v1,12(v0) |
| 116c: | 8fc2001c | lw | v0,28(s8) |
| 1170: | 00021080 | sll | v0,v0,0x2 |
| 1174: | 27c40018 | addiu | a0,s8,24 |
| 1178: | 00821021 | addu | v0,a0,v0 |
| 117c: | ac43000c | sw | v1,12(v0) |
| 1180: | 8fc20018 | lw | v0,24(s8) |
| 1184: | 00021080 | sll | v0,v0,0x2 |
| 1188: | 27c30018 | addiu | v1,s8,24 |
| 118c: | 00621021 | addu | v0,v1,v0 |

```

1190:      8fc30020      lw      v1,32(s8)
1194:      ac43000c      sw      v1,12(v0)
1198:      8fc20018      lw      v0,24(s8)
119c:      24420001      addiu   v0,v0,1
11a0:      afc20018      sw      v0,24(s8)
11a4:      8fc20018      lw      v0,24(s8)
11a8:      28420008      slti   v0,v0,8
11ac:      1440ffd6      bnez   v0,1108 <main+0x70>
11b0:      00000000      nop
11b4:      0000000d      break
11b8:      8fc20024      lw      v0,36(s8)
11bc:      03c0e825      move   sp,s8
11c0:      8fbf004c      lw      ra,76(sp)
11c4:      8fbe0048      lw      s8,72(sp)
11c8:      27bd0050      addiu   sp,sp,80
11cc:      03e00008      jr      ra
11d0:      00000000      nop

000011d4 <__start>:
11d4:      27bdffe0      addiu   sp,sp,-32
11d8:      afbf001c      sw      ra,28(sp)
11dc:      afbe0018      sw      s8,24(sp)
11e0:      03a0f025      move   s8,sp
11e4:      3c1d0f00      lui    sp,0xf00
11e8:      0c000426      jal     1098 <main>
11ec:      00000000      nop
11f0:      00000000      nop
11f4:      03c0e825      move   sp,s8
11f8:      8fbf001c      lw      ra,28(sp)
11fc:      8fbe0018      lw      s8,24(sp)
1200:      27bd0020      addiu   sp,sp,32
1204:      03e00008      jr      ra
1208:      00000000      nop
120c:      00000000      nop

```

look at the code + .c

```

sjtu@sjtu-OptiPlex-3010:~/Documents/DossierNicolasien/PR2$ mips-linux-gnu-objdump -S
bin/insertion-sort.elf

```

```

bin/insertion-sort.elf:      file format elf32-tradlittlemips

```

```

Disassembly of section .text:

```

```

00001000 <minIndex>:

```

```

#define SIZE 8

```

```

int input[] = {60, 41, 46, 50, 44, 3, 84, 80, 55, 57, 91, 22, 21, 12, 64, 59, 71, 34, 81,

```

```
77, 69, 95, 2, 24, 61, 73, 25, 19, 29, 91, 45, 53, 39, 15, 47, 58, 3, 62, 81, 0, 33, 83, 12,
64, 75, 59, 32, 68, 98, 68, 53, 74, 88, 30, 65, 23, 97, 66, 49, 46, 18, 22, 0, 30, 3, 33,
13, 33, 31, 61, 14, 87, 57, 95, 20, 92, 67, 71, 42, 52, 18, 98, 2, 93, 95, 69, 90, 8, 97,
46, 26, 68, 69, 84, 73, 35, 44, 88, 79, 65};
```

```
int minIndex(int *array, int n)
```

```
{
    1000:      27bdffe8      addiu   sp,sp,-24
    1004:      afbe0014      sw      s8,20(sp)
    1008:      03a0f025      move    s8,sp
    100c:      afc40018      sw      a0,24(s8)
    1010:      afc5001c      sw      a1,28(s8)
        int i,minIdx = 0;
    1014:      afc0000c      sw      zero,12(s8)
        for(i = 0; i < n; i++)
    1018:      afc00008      sw      zero,8(s8)
    101c:      10000013      b      106c <minIndex+0x6c>
    1020:      00000000      nop
        {
            if (array[i] < array[minIdx])
    1024:      8fc20008      lw      v0,8(s8)
    1028:      00021080      sll    v0,v0,0x2
    102c:      8fc30018      lw      v1,24(s8)
    1030:      00621021      addu   v0,v1,v0
    1034:      8c430000      lw      v1,0(v0)
    1038:      8fc2000c      lw      v0,12(s8)
    103c:      00021080      sll    v0,v0,0x2
    1040:      8fc40018      lw      a0,24(s8)
    1044:      00821021      addu   v0,a0,v0
    1048:      8c420000      lw      v0,0(v0)
    104c:      0062102a      slt    v0,v1,v0
    1050:      10400003      beqz   v0,1060 <minIndex+0x60>
    1054:      00000000      nop
        {
```

```
            minIdx = i;
    1058:      8fc20008      lw      v0,8(s8)
    105c:      afc2000c      sw      v0,12(s8)
int input[] = {60, 41, 46, 50, 44, 3, 84, 80, 55, 57, 91, 22, 21, 12, 64, 59, 71, 34, 81,
77, 69, 95, 2, 24, 61, 73, 25, 19, 29, 91, 45, 53, 39, 15, 47, 58, 3, 62, 81, 0, 33, 83, 12,
64, 75, 59, 32, 68, 98, 68, 53, 74, 88, 30, 65, 23, 97, 66, 49, 46, 18, 22, 0, 30, 3, 33,
13, 33, 31, 61, 14, 87, 57, 95, 20, 92, 67, 71, 42, 52, 18, 98, 2, 93, 95, 69, 90, 8, 97,
46, 26, 68, 69, 84, 73, 35, 44, 88, 79, 65};
```

```
int minIndex(int *array, int n)
```

```
{
    int i,minIdx = 0;
    for(i = 0; i < n; i++)
    1060:      8fc20008      lw      v0,8(s8)
    1064:      24420001      addiu   v0,v0,1
```

```

1068:      afc20008      sw      v0,8(s8)
106c:      8fc30008      lw      v1,8(s8)
1070:      8fc2001c      lw      v0,28(s8)
1074:      0062102a      slt     v0,v1,v0
1078:      1440ffea      bnez    v0,1024 <minIndex+0x24>
107c:      00000000      nop
      {
          minIdx = i;
      }
  }

      return minIdx;
1080:      8fc2000c      lw      v0,12(s8)
}
1084:      03c0e825      move    sp,s8
1088:      8fbe0014      lw      s8,20(sp)
108c:      27bd0018      addiu   sp,sp,24
1090:      03e00008      jr      ra
1094:      00000000      nop

00001098 <main>:

int main()
{
    1098:      27bdffb0      addiu   sp,sp,-80 // r29, r29, -80 ?
    109c:      afbf004c      sw      ra,76(sp) // 76(r29), r31
    10a0:      afbe0048      sw      s8,72(sp) // 72(r29), r31
    10a4:      03a0f025      move    s8,sp // or r30, r29, r0
        int buf[SIZE];
        int i;

        for(i = 0; i < SIZE; i++)
    10a8:      afc00018      sw      zero,24(s8) // sw 24(r30), r0
    10ac:      1000000f      b      10ec <main+0x54> // beqz r0, 0x3c, jump
    10b0:      00000000      nop
        {
            buf[i] = input[i];
    10b4:      3c020001      lui     v0,0x1 // v0 = 1 << 16?
    10b8:      8fc30018      lw      v1,24(s8) // v1: input[i], s8:
    10bc:      00031880      sll     v1,v1,0x2 // v1 *= 4
    10c0:      24421210      addiu   v0,v0,4624 // 4624 offset ?
    10c4:      00621021      addu    v0,v1,v0 // v0 +=
    10c8:      8c430000      lw      v1,0(v0) // v1 = v0
    10cc:      8fc20018      lw      v0,24(s8) // v0 = i
    10d0:      00021080      sll     v0,v0,0x2 // v0 = 4i
    10d4:      27c40018      addiu   a0,s8,24 // s8 = a0 + 24
    10d8:      00821021      addu    v0,a0,v0 // v0 = v0 + a0 ?
    10dc:      ac43000c      sw      v1,12(v0) // M[v0 + 12] = v1 -> loading in to memory,
v0 = 0xeffffc8+4k

```

```

int main()
{
    int buf[SIZE];
    int i;

    for(i = 0; i < SIZE; i++)
10e0:      8fc20018      lw      v0,24(s8) // i
10e4:      24420001      addiu   v0,v0,1 // i++
10e8:      afc20018      sw      v0,24(s8)
10ec:      8fc20018      lw      v0,24(s8) // i load from memory?
10f0:      28420008      slti    v0,v0,8 // i < 8 ?
10f4:      1440ffef      bnez    v0,10b4 <main+0x1c>
10f8:      00000000      nop
    {
        buf[i] = input[i];
    }

    for(i = 0; i < SIZE; i++)
10fc:      afc00018      sw      zero,24(s8)
1100:      10000028      b       11a4 <main+0x10c>
1104:      00000000      nop
    {
        int minIdx = i + minIndex(&buf[i], SIZE - i);
1108:      27c30024      addiu   v1,s8,36
110c:      8fc20018      lw      v0,24(s8)
1110:      00021080      sll     v0,v0,0x2
1114:      00622021      addu    a0,v1,v0
1118:      24030008      li      v1,8
111c:      8fc20018      lw      v0,24(s8)
1120:      00621023      subu    v0,v1,v0
1124:      00402825      move    a1,v0
1128:      0c000400      jal     1000 <minIndex>
112c:      00000000      nop
1130:      00401825      move    v1,v0
1134:      8fc20018      lw      v0,24(s8)
1138:      00621021      addu    v0,v1,v0
113c:      afc2001c      sw      v0,28(s8)
        int tmp = buf[minIdx];
1140:      8fc2001c      lw      v0,28(s8)
1144:      00021080      sll     v0,v0,0x2
1148:      27c30018      addiu   v1,s8,24
114c:      00621021      addu    v0,v1,v0
1150:      8c42000c      lw      v0,12(v0)
1154:      afc20020      sw      v0,32(s8)
        buf[minIdx] = buf[i];
1158:      8fc20018      lw      v0,24(s8)
115c:      00021080      sll     v0,v0,0x2
1160:      27c30018      addiu   v1,s8,24
1164:      00621021      addu    v0,v1,v0

```

```

1168:      8c43000c      lw      v1,12(v0)
116c:      8fc2001c      lw      v0,28(s8)
1170:      00021080      sll     v0,v0,0x2
1174:      27c40018      addiu   a0,s8,24
1178:      00821021      addu    v0,a0,v0
117c:      ac43000c      sw      v1,12(v0)
      buf[i] = tmp;
1180:      8fc20018      lw      v0,24(s8)
1184:      00021080      sll     v0,v0,0x2
1188:      27c30018      addiu   v1,s8,24
118c:      00621021      addu    v0,v1,v0
1190:      8fc30020      lw      v1,32(s8)
1194:      ac43000c      sw      v1,12(v0)
      for(i = 0; i < SIZE; i++)
      {
          buf[i] = input[i];
      }

      for(i = 0; i < SIZE; i++)
1198:      8fc20018      lw      v0,24(s8)
119c:      24420001      addiu   v0,v0,1
11a0:      afc20018      sw      v0,24(s8)
11a4:      8fc20018      lw      v0,24(s8)
11a8:      28420008      slti    v0,v0,8
11ac:      1440ffd6      bnez    v0,1108 <main+0x70>
11b0:      00000000      nop
      int tmp = buf[minIdx];
      buf[minIdx] = buf[i];
      buf[i] = tmp;
  }

      asm volatile ("break");
11b4:      0000000d      break
      return buf[0];
11b8:      8fc20024      lw      v0,36(s8)
}
11bc:      03c0e825      move    sp,s8
11c0:      8fbf004c      lw      ra,76(sp)
11c4:      8fbe0048      lw      s8,72(sp)
11c8:      27bd0050      addiu   sp,sp,80
11cc:      03e00008      jr      ra
11d0:      00000000      nop

000011d4 <__start>:

void __start()
{
    11d4:      27bdf0e0      addiu   sp,sp,-32
    11d8:      afbf001c      sw      ra,28(sp)

```



```

11dc:      afbe0018      sw      s8,24(sp)
11e0:      03a0f025      move     s8,sp
asm volatile ("lui $sp, 0xf00");
11e4:      3c1d0f00      lui      sp,0xf00
main();
11e8:      0c000426      jal      1098 <main>
11ec:      00000000      nop
// does not return
}
11f0:      00000000      nop
11f4:      03c0e825      move     sp,s8
11f8:      8fbf001c      lw       ra,28(sp)
11fc:      8fbe0018      lw       s8,24(sp)
1200:      27bd0020      addiu    sp,sp,32
1204:      03e00008      jr       ra
1208:      00000000      nop
120c:      00000000      nop

```

Symbols in

```

sjtu@sjtu-OptiPlex-3010:~/Documents/DossierNicolasien/PR2$ mips-linux-gnu-objdump -t -j.data
-j.text bin/insertion-sort.elf

```

```

bin/insertion-sort.elf:      file format elf32-tradlittlemips

```

SYMBOL TABLE:

```

00001000 l    d  .text 00000000 .text
00011210 l    d  .data 00000000 .data
00011210 g      .data 00000000 _fdata
00011210 g    O  .data 00000190 input
000011d4 g    F  .text 00000038 __start
00001000 g      .text 00000000 _ftext
000113a0 g      .data 00000000 __bss_start
00001098 g    F  .text 0000013c main
000113a0 g      .data 00000000 _edata
000113a0 g      .data 00000000 _end
00001000 g    F  .text 00000098 minIndex
000113a0 g      .data 00000000 _fbss

```

Symbol for INPUT in C code :

```

00011210 l    d  .data 00000000 .data
00011210 g      .data 00000000 _fdata
00011210 g    O  .data 00000190 input

```

Loading address of input into register.

```

for(i = 0; i < SIZE; i++)

```

```

10e0:      8fc20018      lw      v0,24(s8)
10e4:      24420001      addiu   v0,v0,1
10e8:      afc20018      sw      v0,24(s8)
10ec:      8fc20018      lw      v0,24(s8)
10f0:      28420008      slti    v0,v0,8
10f4:      1440ffef      bnez    v0,10b4 <main+0x1c>
10f8:      00000000      nop
      {
          buf[i] = input[i];
      }

```

2.2 openDLX

Checking 0x11210, match against input array

0x3c => 60_(10), 0x29 => 41 ...

determine the address to which the store instruction is writing

0x8fc20018 100011 | 11 110 | 0 0010 | 0000 0000 0001 1000 rt = 2, rs = 30, imm = 24

Values in memory are sirted

9 cycles misprediction

11e8: 0c000426 jal 1098

an unconditional branch jump

Statistics

```

----- SIMULATION STATISTICS -----
Cycles: 9
Executed instructions: 6
Performed fetches: 9
Cache statistics:
Icache: used rpol: DIRECT_MAPPED lines: 32 associativity: 1 line_size: 8 total size: 256
Accesses:9 hits: 3 misses: 6 hit rate: 33,33% loaded words: 12
Dcache: used rpol: DIRECT_MAPPED wpol: WRITE_THROUGH lines: 4 associativity: 1 line_size: 4
total size: 16
Accesses: 0 hits: 0 misses: 0 loaded words: 0
Jumps: 1 (taken: 1, not taken: 0) branches_likely: 0 branches_and_link: 1
Branch Target Buffer (128, D_1BIT): hits: 0 misses: 1
Jumps correctly predicted: 0 mispredicted: 1 misprediction rate: 100%
Number of unique jumps: 1
bpc: 0x000011e8 [104] tgts: [0x00001098] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1
Memory accesses: 2 (reads: 0, writes: 2)
ALU forwarded values: 3 (from execute: 3, memory stage: 0, write back: 0)
BCRTL forwarded values: 0 (from execute: 0, memory stage: 0, write back: 0)
STORE forwarded values: 0 (from execute: 0, memory stage: 0, write back: 0)
Total forwarded values: 3 (from execute: 3, memory stage: 0, write back: 0)

```

After the execution

----- SIMULATION STATISTICS -----

Cycles: 1868

Executed instructions: 1441

Performed fetches: 1868

Cache statistics:

Icache: used rpol: DIRECT_MAPPED lines: 32 associativity: 1 line_size: 8 total size: 256

Accesses:1868 hits: 1558 misses: 310 hit rate: 83,4% loaded words: 620

Dcache: used rpol: DIRECT_MAPPED wpol: WRITE_THROUGH lines: 4 associativity: 1 line_size: 4
total size: 16

Accesses: 500 hits: 220 misses: 280 hit rate: 44% loaded words: 280

Jumps: 125 (taken: 110, not taken: 15) branches_likely: 0 branches_and_link: 9

Branch Target Buffer (128, D_1BIT): hits: 115 misses: 10

Jumps correctly predicted: 88 mispredicted: 37 misprediction rate: 29,6%

Number of unique jumps: 10

bpc: 0x00001078 [120] tgts: [0x00001024] a:44 t/nt: 36/8 mp/cp: 16/28 mp-ratio: 0,36

bpc: 0x00001050 [80] tgts: [0x00001060] a:36 t/nt: 31/5 mp/cp: 11/25 mp-ratio: 0,31

bpc: 0x000010f4 [116] tgts: [0x000010b4] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio: 0,22

bpc: 0x000011ac [44] tgts: [0x00001108] a:9 t/nt: 8/1 mp/cp: 2/7 mp-ratio: 0,22

bpc: 0x0000101c [28] tgts: [0x0000106c] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio: 0,12

bpc: 0x00001090 [16] tgts: [0x00001130] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio: 0,12

bpc: 0x00001128 [40] tgts: [0x00001000] a:8 t/nt: 8/0 mp/cp: 1/7 mp-ratio: 0,12

bpc: 0x000010ac [44] tgts: [0x000010ec] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1

bpc: 0x00001100 [0] tgts: [0x000011a4] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1

bpc: 0x000011e8 [104] tgts: [0x00001098] a:1 t/nt: 1/0 mp/cp: 1/0 mp-ratio: 1

Memory accesses: 643 (reads: 500, writes: 143)

ALU forwarded values: 994 (from execute: 957, memory stage: 37, write back: 0)

BCRTL forwarded values: 95 (from execute: 88, memory stage: 7, write back: 0)

STORE forwarded values: 102 (from execute: 101, memory stage: 1, write back: 0)

Total forwarded values: 1191 (from execute: 1146, memory stage: 45, write back: 0)

----- SIMULATION STATISTICS -----